

✓ Libraries Importation

```
from __future__ import print_function, division
from torchvision.datasets import EuroSAT
import torch
import torch.nn as nn
import torch.optim as optim
from torch.optim import lr_scheduler
import numpy as np
import torchvision
from torchvision import datasets, models, transforms
import matplotlib.pyplot as plt
import time
import os
import copy
import random
from torch.utils.data import DataLoader
from itertools import product
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
from torchcam.methods import SmoothGradCAMpp
from torchcam.utils import overlay_mask
from torchvision.transforms.functional import to_pil_image
```

✓ Data Preprocessing

```
# Set random seeds for reproducibility
```

```
random.seed(42)
```

```
torch.manual_seed(42)
```

```
↳ <torch._C.Generator at 0x7d0a5c4c68d0>
```

```
# Data augmentation and normalization for training
```

```
# Just normalization for validation and test
```

```
data_transforms = {
    'train': transforms.Compose([
        transforms.RandomResizedCrop(224, scale=(0.8, 1.0)),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406],
                              [0.229, 0.224, 0.225])
    ]),
    'val': transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406],
                              [0.229, 0.224, 0.225])
    ]),
    'test': transforms.Compose([
        transforms.Resize(256),
        transforms.CenterCrop(224),
        transforms.ToTensor(),
        transforms.Normalize([0.485, 0.456, 0.406],
                              [0.229, 0.224, 0.225])
    ]),
}
```

```
# Can be used to wrap subsets and apply specific transforms to different subsets
```

```
class TransformSubset(torch.utils.data.Dataset):
```

```
    def __init__(self, subset, transform=None):
```

```
        self.subset = subset
```

```
        self.transform = transform
```

```
    def __getitem__(self, index):
```

```
        x, y = self.subset[index]
```

```
        if self.transform:
```

```
            x = self.transform(x)
```

```
        return x, y
```

```

def __len__(self):
    return len(self.subset)

# Download the EuroSAT dataset.
data_dir = './data/eurosat'
eurosat_dataset = EuroSAT(root=data_dir, download=True)

↗ 100%|██████████| 94.3M/94.3M [00:00<00:00, 357MB/s]

# Split dataset into train (49%), val (21%), and test (30%)
# train:val = 70:30
dataset_size = len(eurosat_dataset)
train_size = int(0.49 * dataset_size)
val_size = int(0.21 * dataset_size)
test_size = dataset_size - train_size - val_size

train_subset, val_subset, test_subset = torch.utils.data.random_split(eurosat_dataset, [train_size, val_size, test_size], generator=torch.Ge

# Wrap each subset with its corresponding transform
train_dataset = TransformSubset(train_subset, transform=data_transforms['train'])
val_dataset = TransformSubset(val_subset, transform=data_transforms['val'])
test_dataset = TransformSubset(test_subset, transform=data_transforms['test'])

# Create dataloaders
batch_size = 128
num_workers = 4
dataloaders = {
    'train': DataLoader(train_dataset, batch_size=batch_size, shuffle=True, num_workers=num_workers),
    'val': DataLoader(val_dataset, batch_size=batch_size, shuffle=False, num_workers=num_workers),
    'test': DataLoader(test_dataset, batch_size=batch_size, shuffle=False, num_workers=num_workers)
}
dataset_sizes = {'train': len(train_dataset), 'val': len(val_dataset), 'test': len(test_dataset)}
class_names = eurosat_dataset.class_names

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

class_names
↗ ['AnnualCrop',
   'Forest',
   'HerbaceousVegetation',
   'Highway',
   'Industrial',
   'Pasture',
   'PermanentCrop',
   'Residential',
   'River',
   'SeaLake']

```

✓ Visualization before going into training

```

# Can be used to visualize a batch of images
def imshow(inp, title=None):
    """Imshow for Tensor."""
    inp = inp.numpy().transpose((1, 2, 0))
    mean = np.array([0.485, 0.456, 0.406])
    std = np.array([0.229, 0.224, 0.225])
    inp = std * inp + mean
    inp = np.clip(inp, 0, 1)
    plt.imshow(inp)
    plt.pause(0.001)

# Visualize a batch of training images
inputs, labels = next(iter(dataloaders['train']))
out = torchvision.utils.make_grid(inputs)
imshow(out, title=[class_names[x] for x in labels])

```



✓ Model Selection and Architecture

```
# A custom CNN architecture
class CustomCNN(nn.Module):
    def __init__(self, dropout_rate=0.5):
        super(CustomCNN, self).__init__()
        # 3 convolutional blocks.
        self.features = nn.Sequential(
            # (3x224x224) to (32x112x112)
            nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),
            # (32x112x112) to (64x56x56)
            nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),
            # (64x56x56) to (128x28x28)
            nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2)
        )
        # Classifier with dropout regularization.
        self.classifier = nn.Sequential(
            nn.Dropout(0.5), # Dropout regularizes by preventing overfitting.
            nn.Linear(128 * 28 * 28, 512),
            nn.ReLU(inplace=True),
            nn.Dropout(0.5),
            nn.Linear(512, 10)
        )

    def forward(self, x):
        x = self.features(x)
        # Flatten the feature maps into a vector for the fully connected layers
        x = x.view(x.size(0), -1)
        x = self.classifier(x)
        return x
```

✓ Model Training

```
def train_model(model, criterion, optimizer, scheduler, num_epochs=25, patience=5):

    since = time.time()
    best_model_wts = copy.deepcopy(model.state_dict())
    best_val_loss = float('inf')
    patience_counter = 0
```

```

for epoch in range(num_epochs):
    print('Epoch {}/{}'.format(epoch, num_epochs - 1))
    print('-' * 10)

    for phase in ['train', 'val']:
        if phase == 'train':
            scheduler.step()
            model.train()
        else:
            model.eval()

    running_loss = 0.0
    running_corrects = 0

    for inputs, labels in dataloaders[phase]:
        inputs = inputs.to(device)
        labels = labels.to(device)

        optimizer.zero_grad()

        with torch.set_grad_enabled(phase == 'train'):
            outputs = model(inputs)
            _, preds = torch.max(outputs, 1)
            loss = criterion(outputs, labels)

            if phase == 'train':
                loss.backward()
                optimizer.step()

        running_loss += loss.item() * inputs.size(0)
        running_corrects += torch.sum(preds == labels.data)

    epoch_loss = running_loss / dataset_sizes[phase]
    epoch_acc = running_corrects.double() / dataset_sizes[phase]

    # Monitor training progress, including loss and accuracy
    print('{} Loss: {:.4f} Acc: {:.4f}'.format(phase, epoch_loss, epoch_acc))

    # EarlyStopping check
    if phase == 'val':
        if epoch_loss < best_val_loss:
            best_val_loss = epoch_loss # Update best loss
            best_model_wts = copy.deepcopy(model.state_dict()) # Update best model
            patience_counter = 0 # Reset the counter because loss improves
        else:
            patience_counter += 1 # Add on the counter because loss does not improve
            if patience_counter >= patience: # If loss does not improve within 5 epochs
                print('Early stopping triggered after {} epochs with no improvement in validation loss.'.format(patience))
                model.load_state_dict(best_model_wts)
                time_elapsed = time.time() - since
                print('Training complete in {:.0f}m {:.0f}s'.format(time_elapsed // 60, time_elapsed % 60))
                print('Best val Loss: {:.4f}'.format(best_val_loss))
                return model # Termination

    print()

    time_elapsed = time.time() - since
    print('Training complete in {:.0f}m {:.0f}s'.format(time_elapsed // 60, time_elapsed % 60))
    print('Best val Loss: {:.4f}'.format(best_val_loss))
    model.load_state_dict(best_model_wts)
    return model

```

▼ Train the CNN Model

```

model_cnn = CustomCNN()
model_cnn = model_cnn.to(device)

# Set up loss function, optimizer, and learning rate scheduler.
criterion = nn.CrossEntropyLoss()
optimizer_cnn = optim.SGD(model_cnn.parameters(), lr=0.01, momentum=0.9)
scheduler_cnn = lr_scheduler.StepLR(optimizer_cnn, step_size=7, gamma=0.1)

# Train the CNN model
num_epochs = 25

```

```
model_cnn = train_model(model_cnn, criterion, optimizer_cnn, scheduler_cnn, num_epochs=num_epochs, patience=5)
```

```
Epoch 14/24
-----
train Loss: 0.8466 Acc: 0.7139
val Loss: 0.6784 Acc: 0.7820

Epoch 15/24
-----
train Loss: 0.8419 Acc: 0.7109
val Loss: 0.6782 Acc: 0.7820

Epoch 16/24
-----
train Loss: 0.8406 Acc: 0.7148
val Loss: 0.6840 Acc: 0.7792

Epoch 17/24
-----
train Loss: 0.8244 Acc: 0.7156
val Loss: 0.6746 Acc: 0.7832

Epoch 18/24
-----
train Loss: 0.8241 Acc: 0.7166
val Loss: 0.6664 Acc: 0.7869

Epoch 19/24
-----
train Loss: 0.8324 Acc: 0.7135
val Loss: 0.6524 Acc: 0.7885

Epoch 20/24
-----
train Loss: 0.8117 Acc: 0.7184
val Loss: 0.6710 Acc: 0.7868

Epoch 21/24
-----
train Loss: 0.8210 Acc: 0.7178
val Loss: 0.6680 Acc: 0.7852

Epoch 22/24
-----
train Loss: 0.8211 Acc: 0.7172
val Loss: 0.6865 Acc: 0.7785

Epoch 23/24
-----
train Loss: 0.8308 Acc: 0.7147
val Loss: 0.6631 Acc: 0.7854

Epoch 24/24
-----
train Loss: 0.8321 Acc: 0.7184
val Loss: 0.6688 Acc: 0.7862
Early stopping triggered after 5 epochs with no improvement in validation loss.
Training complete in 5m 55s
Best val Loss: 0.6524
```

▼ Train a pre-trained ResNet18 model

```
# Load a pre-trained ResNet18 and modify the final layer for 10 classes
```

```
model_ft = models.resnet18(pretrained=True)
```

```
num_fts = model_ft.fc.in_features
```

```
model_ft.fc = nn.Linear(num_fts, 10)
```

```
model_ft = model_ft.to(device)
```

```
# Define loss function, optimizer, and learning rate scheduler
```

```
criterion = nn.CrossEntropyLoss()
```

```
optimizer_ft = optim.SGD(model_ft.parameters(), lr=0.001, momentum=0.9)
```

```
exp_lr_scheduler = lr_scheduler.StepLR(optimizer_ft, step_size=7, gamma=0.1)
```

```
# Train the model
```

```
num_epochs = 25
```

```
model_ft = train_model(model_ft, criterion, optimizer_ft, exp_lr_scheduler, num_epochs=num_epochs, patience=5)
```

```
Epoch 25/25
-----
train Loss: 0.8306 Acc: 0.7147
val Loss: 0.6706 Acc: 0.7873
```

```
val Loss: 0.0790 Acc: 0.9743
```

```
Epoch 9/24
```

```
-----
```

```
train Loss: 0.0705 Acc: 0.9794
```

```
val Loss: 0.0779 Acc: 0.9743
```

```
Epoch 10/24
```

```
-----
```

```
train Loss: 0.0693 Acc: 0.9796
```

```
val Loss: 0.0770 Acc: 0.9743
```

```
Epoch 11/24
```

```
-----
```

```
train Loss: 0.0680 Acc: 0.9803
```

```
val Loss: 0.0768 Acc: 0.9758
```

```
Epoch 12/24
```

```
-----
```

```
train Loss: 0.0656 Acc: 0.9804
```

```
val Loss: 0.0765 Acc: 0.9741
```

```
Epoch 13/24
```

```
-----
```

```
train Loss: 0.0652 Acc: 0.9794
```

```
val Loss: 0.0751 Acc: 0.9757
```

```
Epoch 14/24
```

```
-----
```

```
train Loss: 0.0651 Acc: 0.9811
```

```
val Loss: 0.0752 Acc: 0.9746
```

```
Epoch 15/24
```

```
-----
```

```
train Loss: 0.0646 Acc: 0.9811
```

```
val Loss: 0.0782 Acc: 0.9741
```

```
Epoch 16/24
```

```
-----
```

```
train Loss: 0.0633 Acc: 0.9805
```

```
val Loss: 0.0782 Acc: 0.9748
```

```
Epoch 17/24
```

```
-----
```

```
train Loss: 0.0651 Acc: 0.9811
```

```
val Loss: 0.0793 Acc: 0.9746
```

```
Epoch 18/24
```

```
-----
```

```
train Loss: 0.0662 Acc: 0.9803
```

```
val Loss: 0.0773 Acc: 0.9755
```

```
Early stopping triggered after 5 epochs with no improvement in validation loss.
```

```
Training complete in 4m 8s
```

```
Best val Loss: 0.0751
```

✓ Hyperparameter Tuning

```
# Can be used to obtain the accuracy
def evaluate_model(model, dataloader):
    model.eval()
    correct = 0
    total = 0
    with torch.no_grad():
        for inputs, labels in dataloader:
            inputs = inputs.to(device)
            labels = labels.to(device)
            outputs = model(inputs)
            _, preds = torch.max(outputs, 1)
            total += labels.size(0)
            correct += torch.sum(preds == labels.data)
    return correct.double() / total
```

```
# Set parameter ranges for Grid Search
learning_rates = [0.001, 0.01]
optimizers = ['SGD', 'Adam']
dropout_rates = [0.3, 0.5]
```

```
num_epochs = 25
```

```

patience = 5

# Keep a record of the hyperparameters used and their impact on the model
experiment_log = []

for lr_val, opt_name, dropout_val in product(learning_rates, optimizers, dropout_rates):
    print("=====")
    print("Running experiment with:")
    print("Learning Rate: {}".format(lr_val))
    print("Optimizer: {}".format(opt_name))
    print("Dropout Rate: {}".format(dropout_val))

    # Initiate the model with the current dropout rate
    model = CustomCNN(dropout_rate=dropout_val).to(device)

    criterion = nn.CrossEntropyLoss()

    # Set up the optimizer
    if opt_name == 'SGD':
        optimizer = optim.SGD(model.parameters(), lr=lr_val, momentum=0.9)
    elif opt_name == 'Adam':
        optimizer = optim.Adam(model.parameters(), lr=lr_val)

    scheduler = lr_scheduler.StepLR(optimizer, step_size=7, gamma=0.1)

    model = train_model(model, criterion, optimizer, scheduler, num_epochs=num_epochs, patience=patience)

    # Evaluate on the validation set
    val_acc = evaluate_model(model, dataloaders['val'])
    print("Final Validation Accuracy: {:.4f}".format(val_acc))

    # Keep the records
    experiment_log.append({
        'learning_rate': lr_val,
        'optimizer': opt_name,
        'dropout_rate': dropout_val,
        'val_accuracy': val_acc.item()
    })
    print("=====\n")

# Display all records
print("----- Experiment Log -----")
for entry in experiment_log:
    print("LR: {learning_rate}, Optimizer: {optimizer}, Dropout: {dropout_rate} --> Val Acc: {val_accuracy:.4f}"
          .format(**entry))

# Determine the best configuration
best_experiment = max(experiment_log, key=lambda x: x['val_accuracy'])
print("\nBest Hyperparameters:")
print(best_experiment)

```



```
train Loss: 1.7147 Acc: 0.3598
val Loss: 1.5544 Acc: 0.4520
```

Epoch 23/24

```
train Loss: 1.7150 Acc: 0.3613
val Loss: 1.5662 Acc: 0.4460
```

Epoch 24/24

```
train Loss: 1.7253 Acc: 0.3587
val Loss: 1.5606 Acc: 0.4455
```

Training complete in 5m 58s

Best val Loss: 1.5544

Final Validation Accuracy: 0.4520

=====

----- Experiment Log -----

```
LR: 0.001, Optimizer: SGD, Dropout: 0.3 --> Val Acc: 0.8353
LR: 0.001, Optimizer: SGD, Dropout: 0.5 --> Val Acc: 0.8383
LR: 0.001, Optimizer: Adam, Dropout: 0.3 --> Val Acc: 0.7924
LR: 0.001, Optimizer: Adam, Dropout: 0.5 --> Val Acc: 0.7672
LR: 0.01, Optimizer: SGD, Dropout: 0.3 --> Val Acc: 0.7788
LR: 0.01, Optimizer: SGD, Dropout: 0.5 --> Val Acc: 0.7326
LR: 0.01, Optimizer: Adam, Dropout: 0.3 --> Val Acc: 0.1079
LR: 0.01, Optimizer: Adam, Dropout: 0.5 --> Val Acc: 0.4520
```

Best Hyperparameters:

```
{'learning rate': 0.001, 'optimizer': 'SGD', 'dropout rate': 0.5, 'val accuracy': 0.8382716049382717}
```

Train the CNN with the best hyperparameters

```
model_cnn = CustomCNN()
```

```
model_cnn = model_cnn.to(device)
```

```
criterion = nn.CrossEntropyLoss()
```

```
optimizer_cnn = optim.SGD(model_cnn.parameters(), lr=0.001, momentum=0.9)
```

```
scheduler_cnn = lr_scheduler.StepLR(optimizer_cnn, step_size=7, gamma=0.1)
```

```
num_epochs = 25
```

```
model_cnn = train_model(model_cnn, criterion, optimizer_cnn, scheduler_cnn, num_epochs=num_epochs, patience=5)
```




```
train Loss: 0.4954 Acc: 0.8274
val Loss: 0.4821 Acc: 0.8325
```

Epoch 18/24

```
train Loss: 0.4954 Acc: 0.8274
val Loss: 0.4793 Acc: 0.8344
```

Epoch 19/24

```
train Loss: 0.4933 Acc: 0.8275
val Loss: 0.4760 Acc: 0.8363
```

Early stopping triggered after 5 epochs with no improvement in validation loss.

Training complete in 4m 45s

Best val Loss: 0.4695

✓ Evaluation

✓ Evaluate the CNN model

```
model_cnn.eval()
```

```
all_labels_cnn = []
```

```
all_preds_cnn = []
```

```
with torch.no_grad():
```

```
    for inputs, labels in dataloaders['val']:
```

```
        inputs = inputs.to(device)
```

```
        outputs = model_cnn(inputs)
```

```
        _, preds = torch.max(outputs, 1)
```

```
        all_labels_cnn.extend(labels.cpu().numpy())
```

```
        all_preds_cnn.extend(preds.cpu().numpy())
```

```
# Compute overall accuracy for CNN
```

```
acc_cnn = np.mean(np.array(all_labels_cnn) == np.array(all_preds_cnn))
```

```
print("CNN Validation Accuracy: {:.4f}".format(acc_cnn))
```

```
print("\nClassification Report for CNN:")
```

```
print(classification_report(all_labels_cnn, all_preds_cnn, target_names=class_names))
```

```
# Plot confusion matrix for CNN
```

```
cm_cnn = confusion_matrix(all_labels_cnn, all_preds_cnn)
```

```
plt.figure(figsize=(8, 6))
```

```
sns.heatmap(cm_cnn, annot=True, fmt="d", cmap="Blues",
```

```
            xticklabels=class_names, yticklabels=class_names)
```

```
plt.xlabel("Predicted")
```

```
plt.ylabel("True")
```

```
plt.title("CNN - Confusion Matrix")
```


```
plt.show()
```

Custom CNN Validation Accuracy: 0.8365

Classification Report for Custom CNN:

	precision	recall	f1-score	support
AnnualCrop	0.83	0.89	0.86	609
Forest	0.87	0.92	0.89	651
HerbaceousVegetation	0.76	0.83	0.80	611
Highway	0.75	0.62	0.68	544
Industrial	0.90	0.95	0.92	537
Pasture	0.79	0.80	0.79	433
PermanentCrop	0.77	0.70	0.73	527
Residential	0.96	0.91	0.93	623
River	0.77	0.79	0.78	523
SeaLake	0.92	0.91	0.91	612
accuracy			0.84	5670
macro avg	0.83	0.83	0.83	5670
weighted avg	0.84	0.84	0.83	5670

Custom CNN - Confusion Matrix



```
def visualize_misclassifications(model, dataloader, num_images=16):
    model.eval()
    misclassified_imgs = []
    misclassified_labels = []
    misclassified_preds = []

    with torch.no_grad():
        for inputs, labels in dataloader:
            inputs = inputs.to(device)
            outputs = model(inputs)
            _, preds = torch.max(outputs, 1)
            # Identify misclassifications
            for i in range(len(labels)):
                if preds[i].cpu() != labels[i]:
                    misclassified_imgs.append(inputs[i].cpu())
                    misclassified_labels.append(labels[i].cpu().item())
                    misclassified_preds.append(preds[i].cpu().item())
                if len(misclassified_imgs) == num_images:
                    break
            if len(misclassified_imgs) == num_images:
                break
```


break

```

# Plot misclassified images
fig, axes = plt.subplots(int(num_images**0.5), int(num_images**0.5), figsize=(8, 8))
for idx, ax in enumerate(axes.flatten()):
    img = misclassified_imgs[idx].numpy().transpose((1,2,0))
    # Denormalize
    mean = np.array([0.485, 0.456, 0.406])
    std = np.array([0.229, 0.224, 0.225])
    img = std * img + mean
    img = np.clip(img, 0, 1)
    ax.imshow(img)
    ax.axis('off')
    ax.set_title(f"True: {class_names[misclassified_labels[idx]]}\nPred: {class_names[misclassified_preds[idx]]}")
plt.tight_layout()
plt.show()

```

```
visualize_misclassifications(model_cnn, dataloaders['test'], num_images=9)
```



▼ Evaluate the pre-trained ResNet18 model

```

model_ft.eval()

all_labels_ft = []
all_preds_ft = []

with torch.no_grad():
    for inputs, labels in dataloaders['val']:
        inputs = inputs.to(device)
        outputs = model_ft(inputs)
        _, preds = torch.max(outputs, 1)
        all_labels_ft.extend(labels.cpu().numpy())

```

```
all_preds_ft.extend(preds.cpu().numpy())


# Compute overall accuracy for ResNet18
acc_ft = np.mean(np.array(all_labels_ft) == np.array(all_preds_ft))
print("ResNet18 Validation Accuracy: {:.4f}".format(acc_ft))
print("\nClassification Report for ResNet18:")
print(classification_report(all_labels_ft, all_preds_ft, target_names=class_names))

# Plot confusion matrix for ResNet18
cm_ft = confusion_matrix(all_labels_ft, all_preds_ft)
plt.figure(figsize=(8, 6))
sns.heatmap(cm_ft, annot=True, fmt="d", cmap="Blues",
            xticklabels=class_names, yticklabels=class_names)
plt.xlabel("Predicted")
plt.ylabel("True")
plt.title("ResNet18 - Confusion Matrix")
plt.show()
```

🔗 ResNet18 Validation Accuracy: 0.9757

Classification Report for ResNet18:


	precision	recall	f1-score	support
AnnualCrop	0.95	0.99	0.97	609
Forest	0.99	1.00	0.99	651
HerbaceousVegetation	0.95	0.97	0.96	611
Highway	0.97	0.98	0.98	544
Industrial	0.99	0.99	0.99	537
Pasture	0.97	0.95	0.96	433
PermanentCrop	0.97	0.92	0.95	527
Residential	1.00	0.99	1.00	623
River	0.97	0.97	0.97	523
SeaLake	0.99	1.00	0.99	612
accuracy			0.98	5670
macro avg	0.98	0.97	0.97	5670
weighted avg	0.98	0.98	0.98	5670




```
visualize_misclassifications(model_ft, dataloaders['test'], num_images=9)
```




True: PermanentCrop
Pred: HerbaceousVegetation




True: Forest
Pred: Pasture




True: HerbaceousVegetation
Pred: Highway



True: River
Pred: SeaLake




True: HerbaceousVegetation
Pred: PermanentCrop



True: River
Pred: AnnualCrop




True: PermanentCrop
Pred: Pasture



True: Residential
Pred: HerbaceousVegetation



True: AnnualCrop
Pred: SeaLake



✓ Fine-Tuning

Build a more sophisticated CNN architecture

```
class FineTunedCNN(nn.Module):
    def __init__(self, dropout_rate=0.5):
        super(FineTunedCNN, self).__init__()
        self.features = nn.Sequential(
            # 2 conv layers + pooling
            nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(inplace=True),
            nn.Conv2d(32, 32, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),

            # 2 conv layers + pooling
            nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
            nn.Conv2d(64, 64, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(64),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),

            # 2 conv layers + pooling
            nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(inplace=True),
            nn.Conv2d(128, 128, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(128),
```

```

        nn.ReLU(inplace=True),
        nn.MaxPool2d(kernel_size=2, stride=2),

        # 2 conv layers + pooling
        nn.Conv2d(128, 256, kernel_size=3, stride=1, padding=1),
        nn.BatchNorm2d(256),
        nn.ReLU(inplace=True),
        nn.Conv2d(256, 256, kernel_size=3, stride=1, padding=1),
        nn.BatchNorm2d(256),
        nn.ReLU(inplace=True),
        nn.MaxPool2d(kernel_size=2, stride=2)
    )

    self.classifier = nn.Sequential(
        nn.Dropout(dropout_rate),
        nn.Linear(256 * 14 * 14, 512),
        nn.ReLU(inplace=True),
        nn.Dropout(dropout_rate),
        nn.Linear(512, 10)
    )

    def forward(self, x):
        x = self.features(x)
        x = x.view(x.size(0), -1)
        x = self.classifier(x)
        return x

model_fine_cnn = FineTunedCNN()
model_fine_cnn = model_fine_cnn.to(device)

criterion = nn.CrossEntropyLoss()
optimizer_fine_cnn = optim.SGD(model_fine_cnn.parameters(), lr=0.001, momentum=0.9)
scheduler_fine_cnn = lr_scheduler.StepLR(optimizer_fine_cnn, step_size=7, gamma=0.1)

# Train the new CNN model with the best hyperparameters defined in the previous section
num_epochs = 25
model_fine_cnn = train_model(model_fine_cnn, criterion, optimizer_fine_cnn, scheduler_fine_cnn, num_epochs=num_epochs, patience=5)

```



```
train Loss: 0.3121 Acc: 0.8990
val Loss: 0.2926 Acc: 0.8995
```

Epoch 22/24

```
train Loss: 0.3128 Acc: 0.8930
val Loss: 0.2940 Acc: 0.9000
```

Epoch 23/24

```
train Loss: 0.3135 Acc: 0.8922
val Loss: 0.2940 Acc: 0.8981
```

Early stopping triggered after 5 epochs with no improvement in validation loss.

Training complete in 8m 34s

Best val Loss: 0.2874

```
model_fine_cnn.eval()
```

```
all_labels_cnn = []
```

```
all_preds_cnn = []
```

```
with torch.no_grad():
```

```
    for inputs, labels in dataloaders['val']:
```

```
        inputs = inputs.to(device)
```

```
        outputs = model_fine_cnn(inputs)
```

```
        _, preds = torch.max(outputs, 1)
```

```
        all_labels_cnn.extend(labels.cpu().numpy())
```

```
        all_preds_cnn.extend(preds.cpu().numpy())
```

```
acc_cnn = np.mean(np.array(all_labels_cnn) == np.array(all_preds_cnn))
```

```
print("Fine-tuned CNN Validation Accuracy: {:.4f}".format(acc_cnn))
```

```
print("\nClassification Report for Fine-tuned CNN:")
```

```
print(classification_report(all_labels_cnn, all_preds_cnn, target_names=class_names))
```

```
cm_cnn = confusion_matrix(all_labels_cnn, all_preds_cnn)
```

```
plt.figure(figsize=(8, 6))
```

```
sns.heatmap(cm_cnn, annot=True, fmt="d", cmap="Blues",
```


```
            xticklabels=class_names, yticklabels=class_names)
```

```
plt.xlabel("Predicted")
```

```
plt.ylabel("True")
```


```
plt.title("Fine-tuned CNN - Confusion Matrix")
```

```
plt.show()
```

 Fine-tuned CNN Validation Accuracy: 0.9021

Classification Report for Fine-tuned CNN:

	precision	recall	f1-score	support
AnnualCrop	0.88	0.91	0.90	609
Forest	0.96	0.99	0.98	651
HerbaceousVegetation	0.85	0.86	0.86	611
Highway	0.81	0.76	0.79	544
Industrial	0.92	0.96	0.94	537
Pasture	0.88	0.87	0.88	433
PermanentCrop	0.84	0.81	0.82	527
Residential	0.99	0.96	0.97	623
River	0.84	0.85	0.85	523
SeaLake	0.99	1.00	0.99	612
accuracy			0.90	5670
macro avg	0.90	0.90	0.90	5670
weighted avg	0.90	0.90	0.90	5670



```
visualize_misclassifications(model_fine_cnn, dataloaders['test'], num_images=9)
```




✓ GradCam to understand what regions the model is paying attention to


```
# Untransformed validation dataset
val_dataset_raw = TransformSubset(val_subset, transform=None)

# Get the first image and its label from the raw validation dataset.
original_img, original_label = val_dataset_raw[0]

# Display the original image and its label
plt.figure(figsize=(6, 6))
plt.imshow(original_img)
plt.title(f"Original Validation Image - Label: {original_label}")
plt.axis("off")
plt.show()
```



Original Validation Image - Label: 8



```

preprocess = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406],
                          [0.229, 0.224, 0.225])
])

# Preprocess the image
input_tensor = preprocess(original_img).unsqueeze(0)

input_tensor.shape

torch.Size([1, 3, 224, 224])

with SmoothGradCAMpp(model_fine_cnn, target_layer="features.8") as cam_extractor:
    input_tensor = input_tensor.to(device)
    out = model_fine_cnn(input_tensor)
    pred_class = out.squeeze(0).argmax().item()
    activation_map = cam_extractor(pred_class, out)


result = overlay_mask(original_img,
                       to_pil_image(activation_map[0].squeeze(0), mode='F'), alpha=0.5)

plt.imshow(result)
plt.axis('off')
plt.title(f"GradCAM - CNN - Predicted: {pred_class}")
plt.show()

```



GradCAM - CNN - Predicted: 8



```

with SmoothGradCAMpp(model_ft, target_layer="layer4") as cam_extractor:
    input_tensor = input_tensor.to(device)
    out = model_ft(input_tensor)
    pred_class = out.squeeze(0).argmax().item()
    activation_map = cam_extractor(pred_class, out)


result = overlay_mask(original_img,
                       to_pil_image(activation_map[0].squeeze(0), mode='F'), alpha=0.5)

plt.imshow(result)
plt.axis('off')
plt.title(f"GradCAM - ResNet18 - Predicted: {pred_class}")
plt.show()

```



GradCAM - ResNet18 - Predicted: 8



✓ Final Model Testing

```

test_acc = evaluate_model(model_fine_cnn, dataloaders['test'])
print('Test Accuracy for the fine-tuned CNN model: {:.4f}'.format(test_acc))

```



Test Accuracy for the fine-tuned CNN model: 0.8995

```

test_acc = evaluate_model(model_ft, dataloaders['test'])
print('Test Accuracy for the pre-trained ResNet18 model: {:.4f}'.format(test_acc))

```



Test Accuracy for the pre-trained ResNet18 model: 0.9716

```
# make sure everything is printing  
# make sure everything is printing
```